

cross_validation_nested

May 29, 2021

1 Nested cross-validation

In this notebook, we show a pattern called **nested cross-validation** which should be used when you want to both evaluate a model and tune the model's hyperparameters.

Cross-validation is a powerful tool to evaluate the statistical performance of a model. It is also used to select the best model from a pool of models. This pool of models can be the same family of predictor but with different parameters. In this case, we call this procedure **hyperparameter tuning**.

We could also imagine that we would like to choose among heterogeneous models that will similarly use the cross-validation.

Before we go into details regarding the nested cross-validation, we will first recall the pattern used to fine tune a model's hyperparameters.

Let's load the breast cancer dataset.

```
[ ]: from sklearn.datasets import load_breast_cancer  
  
data, target = load_breast_cancer(return_X_y=True)
```

Now, we'll make a minimal example using the utility `GridSearchCV` to find the best parameters via cross-validation.

```
[ ]: from sklearn.model_selection import GridSearchCV  
from sklearn.svm import SVC  
  
param_grid = {"C": [0.1, 1, 10], "gamma": [.01, .1]}  
model_to_tune = SVC()  
  
search = GridSearchCV(estimator=model_to_tune, param_grid=param_grid,  
                      n_jobs=-1)  
search.fit(data, target)
```

We recall that `GridSearchCV` will train a model with some specific parameter on a training set and evaluate it on testing. However, this evaluation is done via cross-validation using the `cv` parameter. This procedure is repeated for all possible combinations of parameters given in `param_grid`.

The attribute `best_params_` will give us the best set of parameters that maximize the mean score on the internal test sets.

```
[ ]: print(f"The best parameter found are: {search.best_params_}")
```

We can now show the mean score obtained using the parameter `best_score_`.

```
[ ]: print(f"The mean score in CV is: {search.best_score_:.3f}")
```

At this stage, one should be extremely careful using this score. The misinterpretation would be the following: since the score was computed on a test set, it could be considered our model's testing score.

However, we should not forget that we used this score to pick-up the best model. It means that we used knowledge from the test set (i.e. test score) to decide our model's training parameter.

Thus, this score is not a reasonable estimate of our testing error. Indeed, we can show that it will be too optimistic in practice. The good way is to use a “nested” cross-validation. We will use an inner cross-validation corresponding to the previous procedure shown to optimize the hyperparameters. We will also include this procedure within an outer cross-validation, which will be used to estimate the testing error of our tuned model.

In this case, our inner cross-validation will always get the training set of the outer cross-validation, making it possible to compute the testing score on a completely independent set.

We will show below how we can create such nested cross-validation and obtain the testing score.

```
[ ]: from sklearn.model_selection import cross_val_score, KFold

# Declare the inner and outer cross-validation
inner_cv = KFold(n_splits=4, shuffle=True, random_state=0)
outer_cv = KFold(n_splits=4, shuffle=True, random_state=0)

# Inner cross-validation for parameter search
model = GridSearchCV(
    estimator=model_to_tune, param_grid=param_grid, cv=inner_cv, n_jobs=-1)

# Outer cross-validation to compute the testing score
test_score = cross_val_score(model, data, target, cv=outer_cv, n_jobs=-1)
print(f"The mean score using nested cross-validation is: "
      f"{test_score.mean():.3f} +/- {test_score.std():.3f}")
```

In the example above, the reported score is more trustful and should be close to production's expected statistical performance.

We will illustrate the difference between the nested and non-nested cross-validation scores to show that the latter one will be too optimistic in practice. In this regard, we will repeat several time the experiment and shuffle the data differently. Besides, we will store the score obtain with and without the nested cross-validation.

```
[ ]: test_score_not_nested = []
test_score_nested = []
```

```
N_TRIALS = 20
```

```

for i in range(N_TRIALS):
    inner_cv = KFold(n_splits=4, shuffle=True, random_state=i)
    outer_cv = KFold(n_splits=4, shuffle=True, random_state=i)

    # Non_nested parameter search and scoring
    model = GridSearchCV(estimator=model_to_tune, param_grid=param_grid,
                          cv=inner_cv, n_jobs=-1)
    model.fit(data, target)
    test_score_not_nested.append(model.best_score_)

    # Nested CV with parameter optimization
    test_score = cross_val_score(model, data, target, cv=outer_cv, n_jobs=-1)
    test_score_nested.append(test_score.mean())

```

We can merge the data together and make a box plot of the two strategies.

```

[ ]: import pandas as pd

all_scores = {
    "Not nested CV": test_score_not_nested,
    "Nested CV": test_score_nested,
}
all_scores = pd.DataFrame(all_scores)

[ ]: import matplotlib.pyplot as plt

color = {"whiskers": "black", "medians": "black", "caps": "black"}
all_scores.plot.box(color=color, vert=False)
plt.xlabel("Accuracy")
_ = plt.title("Comparison of mean accuracy obtained on the test sets with\n"
              "and without nested cross-validation")

```

We observe that the model's statistical performance with the nested cross-validation is not as good as the non-nested cross-validation.

As a conclusion, when optimizing parts of the machine learning pipeline (e.g. hyperparameter, transform, etc.), one needs to use nested cross-validation to evaluate the statistical performance of the predictive model. Otherwise, the results obtained without nested cross-validation are over-optimistic.